

Unit 2.3 Sets and Graphs

Algorithms

EE/NTHU




Mar. 17, 2021

Graphs



- A graph, G , consists of two sets V and E .
 - The set V is a finite, nonempty set of **vertices**.
 - The set E is a set of pairs of vertices; these pairs are called **edges**.
 - They are also denoted by $V(G)$ and $E(G)$.
 - And the graph is also denoted by $G(V, E)$.
- In an **undirected graph** the pair of vertices representing any edge is unordered.
 - Thus, the pairs (u, v) and (v, u) represent the same edge.
- In a **directed graph** each edge is represented by a direct pair $\langle u, v \rangle$; u is the **tail** and v is the **head**.
 - And, $\langle u, v \rangle$ and $\langle v, u \rangle$ represent two different edges.
 - In a directed graph the edges are drawn as arrows and they are drawn from tail to head.

Graph Examples

- Examples



- Note that G_1 and G_2 are undirected graphs; G_3 is a directed graph.
- G_1 is a complete graph; G_2 is a tree.
- The following graphs are **not** studied in our classes:
 - Graphs with self-edges, which have edges connecting the same vertex.
 - Multi-graph, which has multiple edges between the same two vertices.



Adjacency

- The number of distinct unordered pairs (u, v) with $u \neq v$ in a graph with n vertices is $n(n-1)/2$.
 - This is the maximum number of edges in any n -vertex, undirected graph.
 - An n -vertex, undirected graph with exactly $n(n-1)/2$ edges is said to be **complete**.
 - In case of a direct graph with n vertices, the maximum number of edges is $n(n-1)$.
- If (u, v) is an edge in $E(G)$, then we say vertices u and v are **adjacent** and edge (u, v) is **incident** on vertices u and v .
- If $\langle u, v \rangle$ is a directed edge, the vertex u is **adjacent to** v and v is **adjacent from** u .
 - The edge $\langle u, v \rangle$ is **incident** to u and v .
- A **subgraph** of $G = (V, E)$ is a graph $G' = (V', E')$ such that $V'(G') \subseteq V(G)$ and $E'(G') \subseteq E(G)$.

Paths and Cycles

- A **path** from vertex u to vertex v in a graph $G = (V, E)$ is a sequence of vertices $u, i_1, i_2, \dots, i_k, v$, such that $(u, i_1), (i_1, i_2), \dots, (i_k, v)$ are edges in $E(G)$.
 - If $G = (V, E)$ is directed, then the path consists of the edges $\langle u, i_1 \rangle, \langle i_1, i_2 \rangle, \dots, \langle i_k, v \rangle$ in $E(G)$.
- The **length** of a path is the number of edges on it.
- A **simple path** is a path in which all vertices except possibly the first and the last are distinct.
- A **cycle** is a simple path in which the first and the last vertices are the same.
 - A **directed cycle** is a cycle in a directed graph.
- In an undirected graph $G = (V, E)$, two vertices u and v are said to be **connected** if and only if there is a path in G from u to v .
- An undirected graph is said to be **connected** if and only if for every pair of distinct vertices u and v in $V(G)$, there is a path from u to v .
- A **connected component** or simply a component H of an undirected graph is a **maximal** connected subgraph.
 - By maximal we mean that G contains no other subgraph that is both connected and properly contains H .


Graph Degrees

- A **tree** is a connected acyclic (contain no cycles) graph.
- A directed graph $G = (V, E)$ is said to be **strongly connected** if and only if for every pair of distinct vertices u and v in $V(G)$, there is a directed path from u to v and also from v to u .
- A **strongly connected component** is a maximal subgraph that is strongly connected.
- The **degree** of a vertex is the number of edges incident to that vertex.
- If $G = (V, E)$ is a directed graph, we define the **in-degree** of a vertex v to be the number of edges for which v is the head.
 - The **out-degree** is defined to be the number of edges for which v is the tail.
- If d_i is the degree of vertex i in a graph $G = (V, E)$ with n vertices, then the number of edge is


$$e = \left(\sum_{i=1}^n d_i \right) / 2. \quad (2.1)$$

Graph Representation – Adjacency Matrix


- Let $G = (V, E)$ be a graph with n vertices, $n \geq 1$. The **adjacency matrix** A is a two-dimensional $n \times n$ matrix with the property that $A[i, j] = 1$ if and only if the edge (i, j) ($\langle i, j \rangle$ for a directed graph) is in $E(G)$. $A[i, j] = 0$ if there is no such edge in $E(G)$.



$$\begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$



$$\begin{bmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \end{bmatrix}$$



$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

Graph Representation – Adjacency Matrix, II


- The adjacency matrix for an undirected graph is symmetric.
 - This is due to that if the edge (i, j) is in $E(G)$ then the edge (j, i) is also in $E(G)$.
- The adjacency matrix for a directed graph may not be symmetric.
- The space needed to represent for a adjacency matrix is n^2 bits.
 - The undirected graph needs only half of this space.
- For an undirected graph the degree of any vertex i is its row sum:

$$\sum_{j=1}^n A[i][j]. \quad (2.2)$$



- For a directed graph the row sum is the out-degree and the column sum is the in-degree.
- The adjacency matrix approach to represent the graph is not the most efficient way in both space and execution time.
 - It does not take advantage of the sparsity of the graph.
 - For example, the time complexity to find the number of edges of a graph, with n vertices, represented by a adjacency matrix is $\mathcal{O}(n^2)$.

Graph Representation – Adjacency Lists


- A graph, $G = (V, E)$, of n vertices, can also be represented by n linked lists.
 - Each vertex has a linked list to represent the adjacent vertices.
- Examples




list array




list array



Adjacency Lists – Examples



list array




- For an undirected graph with n vertices and e edges, the adjacency list representation requires n head nodes and $2e$ list nodes.
- The degree of any vertex in an undirected graph can be determined by counting the number of nodes in the adjacency list.
 - Hence the total number of edges can be determined in $\mathcal{O}(n + e)$ time.
- For a directed graph, the out-degree of any vertex is again the number of nodes of its adjacency list.
- The in-degree may need to have another **inverse adjacency list**.


- In many applications, the edges of a graph have weights assigned to them. Thus, the adjacency matrix and adjacency lists need to accommodate these weights information.
- The adjacency matrix can store the weight of edge $\langle i, j \rangle$ to $A[i][j]$ directly.
 - No extra storage is required.
 - Space complexity is $\mathcal{O}(n^2)$.
- For the adjacency lists, each node of the list needs to have an additional field to store the weight.
 - In terms of space complexity, it is still the same as $\Theta(e)$, where e is the number of edges in $G = (V, E)$, Or, in worst-case $\mathcal{O}(n^2)$.

Network Connectivity Problem

- Given a network below, is node u connected to node v ?



Network Connectivity Problem, II



- A smaller instance is shown above.
- One solution approach is to form sets of connected nodes, S_i .
- If there is a S_k such that $u, v \in S_k$, then u is connected to v .

Network Connectivity Problem, Algorithm

- A generic algorithm for network connectivity problem is shown below.


Algorithm 2.3.1. Connectivity – Generic.

```
// Given  $G(V, E)$  and  $u, v \in V$ , find if  $u$  and  $v$  are connected.
// Input:  $G, u, v$ 
// Output: true if connected, false otherwise.
1 Algorithm Connected( $G, u, v$ )
2 {
3     for each  $v_i \in V$  do  $S_i := \{v_i\}$  ; // One element for each set.
4     for each  $e = (v_i, v_j)$  do { // Connected vertices
5          $S_i := \text{SetFind}(v_i)$  ;
6          $S_j := \text{SetFind}(v_j)$  ;
7          $S_i := S_i \cup S_j$  ; // Set union.
8     }
9     if  $\text{SetFind}(u) = \text{SetFind}(v)$  then return true ;
10    return false ;
11 }
```

- The time complexity is dominated by the loop on lines 4-8
 - Iterations: $\mathcal{O}(|E|)$.
 - Two **SetFind** and one **SetUnion** per iteration.

Disjoint Sets


- Disjoint sets
 - Assume the elements are numbered $1, 2, \dots, n$.
 - Disjoint sets S_i, S_j such that $S_i \cap S_j = \emptyset, i \neq j$.
 - Forest can be used to represent disjoint sets
- Operations important to set manipulations
 - **Union**: Merge two disjoint sets into one.
 - **Find**(i): Given an element i find the set that contains i .
- Example: $S_1 = \{1, 7, 8, 9\}, S_2 = \{2, 5, 10\}, S_3 = \{3, 4, 6\}$.



- Note that for the forest the link is pointing from child to parent.
 - In this way, the set name can be found by following the pointers.

Disjoint Sets – Array Representation


- Simple array can also be used to represent disjoint sets.
- Example: $S_1 = \{1, 7, 8, 9\}, S_2 = \{2, 5, 10\}, S_3 = \{3, 4, 6\}$.



- A single array, p , can represent the disjoint sets.

Disjoints Sets – Union

- Two disjoint sets can be united easily.
- Example



- Both scenarios are legal and efficient.
- Union of two sets are done by setting one of the roots to be the parent of another root.

Disjoint Sets – Algorithms

- Using the array p to represent the disjoint sets, then the following algorithms perform the desired operations.

Algorithm 2.3.2. Set Union.

```
// Form union of two sets with roots,  $i$  and  $j$ .  
// Input: roots,  $i$  and  $j$   
// Output: none.  
1 Algorithm SetUnion( $i, j$ )  
2 {  
3      $p[i] := j$ ;  
4 }
```

Algorithm 2.3.3. Set Find.

```
// Find the set that element  $i$  is in.  
// Input: element  $i$   
// Output: root element of the set.  
1 Algorithm SetFind( $i$ )  
2 {  
3     while ( $p[i] \geq 0$ ) do  $i := p[i]$ ;  
4     return  $i$ ;  
5 }
```

Disjoint Sets – Weighting Rule

- Algorithm `SetFind(i)` has the complexity $\mathcal{O}(h)$, where h is the height of the tree the element i is in.
- Algorithm `SetUnion(i, j)` has the time complexity $\mathcal{O}(1)$.
 - However, each union operation increases the height of the tree by 1.
 - Thus, after some union operations the tree might become skewed and the execution time of `SetFind` increases.
 - This issue can be alleviated by using the **weighting rule**.

Definition 2.3.4. Weighting rule for set union.

If the number of nodes in the tree with root i is less than the number of nodes in the tree with root j , then make j the parent of i ; otherwise make i the parent of j .

- In order to implement the weighting rule, we need to know the number of elements in each set. This can be done using the root location in the p array. Set it to $-count(i)$, $count(i)$ is the number of elements in set i .
- Example: the disjoint sets can be represented as

	1	2	3	4	5	6	7	8	9	10
p	-4	5	-3	3	-3	3	1	1	1	5

Disjoint Sets – Weighted Set Union

Algorithm 2.3.5. Weighted Set Union.

```
// Form union of two sets with roots,  $i$  and  $j$ , using the weighting rule.
// Input: roots of two sets  $i, j$ 
// Output: none.
1 Algorithm WeightedUnion( $i, j$ )
2 {
3      $temp := p[i] + p[j]$ ; // Note that  $temp < 0$ .
4     if ( $p[i] > p[j]$ ) then { //  $i$  has fewer elements.
5          $p[i] := j$ ;
6          $p[j] := temp$ ;
7     }
8     else { //  $j$  has fewer elements.
9          $p[j] := i$ ;
10         $p[i] := temp$ ;
11    }
12 }
```

- Using this algorithm, the depth of the union tree can be controlled.

Weighted Set Union – Complexity

Lemma 2.3.6.

Assume that we start with a forest of trees, each having one element. Let T be a tree with m nodes created as a result of a sequence of unions each performed using **WeightedUnion** algorithm. The height of T is no greater than $\lfloor \lg m \rfloor + 1$.

Proof. The first step is true when two sets of one element are united. Assume the Lemma is true for the first $m - 1$ operations, consider the last step of the union operations, **WeightedUnion**(k, j). If set j has a elements, then set k has $m - a$ elements. And, $1 \leq a \leq m/2$. The height of T must be the same as that of k or one more than that of j . In the former case, the height of T is $\leq \lfloor \lg(m - a) \rfloor + 1 \leq \lfloor \lg m \rfloor + 1$. In the latter case, the height of T is $\leq \lfloor \lg a \rfloor + 2 \leq \lfloor \lg m/2 \rfloor + 2 \leq \lfloor \lg m \rfloor + 1$. \square

- Thus, the union set created using Algorithm **WeightedUnion** has no more than $\lfloor \lg m \rfloor + 1$ levels.
- And the time complexity of **Find** algorithm on the resulting set is $\mathcal{O}(\lg m)$.

Disjoint Sets – Collapsing Find

- The height of a set may still be improved using the **collapsing rule**.

Definition 2.3.7. Collapsing Rule.

If j is an element on the path from i to its root and $p[i] \neq \text{root}(i)$, then set $p[j]$ to $\text{root}(i)$.

- The **CollapsingFind** algorithm below utilizes this rule.

Algorithm 2.3.8. Collapsing Find.

```
// Find the root of  $i$ , and collapsing the elements on the path.
// Input: an element  $i$ 
// Output: root of the set containing  $i$ .
1 Algorithm CollapsingFind( $i$ )
2 {
3      $r := i$ ; // Initialized  $r$  to  $i$ .
4     while ( $p[r] > 0$ ) do  $r := p[r]$ ; // Find the root.
5     while ( $i \neq r$ ) do { // Collapse the elements on the path.
6          $s := p[i]$ ;  $p[i] := r$ ;  $i := s$ ;
7     }
8     return  $r$ ;
9 }
```

Ackermann's Function

Definition 2.3.9. Ackermann's function.

The Ackermann's function is defined as

$$\begin{aligned} A(1, j) &= 2^j && \text{for } j \geq 1, \\ A(i, 1) &= A(i-1, 2) && \text{for } i \geq 2, \\ A(i, j) &= A(i-1, A(i, j-1)) && \text{for } i, j \geq 2. \end{aligned} \quad (2.3)$$

Also define

$$\alpha(p, q) = \min\{z \geq 1 \mid A(z, \lfloor \frac{p}{q} \rfloor) > \lg q\}, p \geq q \geq 1. \quad (2.4)$$

$$\begin{array}{llll} A(1, 1) = 2 & A(1, 2) = 4 & A(1, 3) = 8 & A(1, 4) = 16 \\ A(2, 1) = 4 & A(2, 2) = 16 & A(2, 3) = 2^{16} & A(2, 4) = 2^{65536} \\ A(3, 1) = 16 & A(3, 2) \gg 2^{65536} & & \\ A(4, 1) \gg 2^{65536} & & & \end{array}$$

- A is very fast growing function and α is a very slow growing function.
- Note that $A(3, 1) = 16$, $\alpha(p, q) \leq 3$ for $q < 2^{16} = 65,536$ and $p > q$.
- Since $A(4, 1)$ is a very large number, $\alpha(p, q) \leq 4$ for all practical purposes.

Tarjan and Van Leeuwen Bound

Lemma 2.3.10. Tarjan and Van Leeuwen bounds.

Assume that we start with a forest of trees, each having one node. Let $T(f, u)$ be the maximum time required to process any intermixed sequence of f finds and u unions. Assume that $u \geq n/2$, then

$$k_1 \left(n + f \cdot \alpha(f + n, n) \right) \leq T(f, u) \leq k_2 \left(n + f \cdot \alpha(f + n, n) \right) \quad (2.5)$$

for some positive constants k_1 and k_2 .

- Proof please see textbook [Cormen], pp. 575-581.
- Thus, manipulating disjoint sets are rather efficient.
- Though algorithms (2.3.2), (2.3.3), (2.3.5), and (2.3.8) assume the disjoint sets are represented using a simple array, they can be implemented if the disjoint sets are represented using linked lists as well.
- The complexities are the same with either data structure.

- Graphs
 - Definitions.
 - Adjacency matrix.
 - Adjacency lists.
- Network connectivity problem
- Disjoint sets.
 - Set union.
 - Set find.
 - Weighted set union.
 - Collapsing set find.

